

```
Last login: Tue Nov 19 22:42:06 on ttys009
> cd ~/OneDrive/2024\ Fall/MLE/workspace/mod3-Navxihziq
> clear
> source .venv/bin/activate
> python project/parallel_check.py
MAP
OMP: Info #276: omp_set_nested routine deprecated, please use omp_set_max_active_levels instead.
```

```
=====
Parallel Accelerator Optimizing: Function tensor_map.<locals>._map,
/Users/qizhixuan/Library/CloudStorage/OneDrive-Personal/2024
Fall/MLE/workspace/mod3-Navxihziq/minitorch/fast_ops.py (163)
=====
```

```
Parallel loop listing for Function tensor_map.<locals>._map, /Users/qizhixuan/Library/CloudStorage/OneDrive-Personal/2024 Fall/MLE/workspace/mod3-Navxihziq/minitorch/fast_ops.py (163)
```

```
----- loop #10 -----
def _map(
    out: Storage,
    out_shape: Shape,
    out_strides: Strides,
    in_storage: Storage,
    in_shape: Shape,
    in_strides: Strides,
) -> None:
    # TODO: Implement for Task 3.1.
    # check if out, in are stride-aligned
    # if out_strides == in_strides:
    #     for i in prange(len(out)):
    #         out[i] = fn(in_storage[i])
    # else:
    # TODO: check if out, in are stride-aligned
    # coerce the shape to int32
    out_shape = out_shape.astype(np.int32)
    in_shape = in_shape.astype(np.int32)
    for i in prange(len(out)):
        out_index = np.zeros(len(out_shape), dtype=np.int32) # buffer----- #2
        in_index = np.zeros(len(in_shape), dtype=np.int32) # buffer----- #0
        to_index(i, out_shape, out_index) #1
        broadcast_index(out_index, out_shape, in_shape, in_index)
        out[i] = fn(in_storage[index_to_position(in_index, in_strides)])
----- Fusing loops -----
```

Attempting fusion of parallel loops (combines loops with similar properties)...
Following the attempted fusion of parallel for-loops there are 3 parallel for-loop(s) (originating from loops labelled: #2, #0, #1).

----- Optimising loop nests -----
Attempting loop nest rewrites (optimising for the largest parallel loops)...

+-2 is a parallel loop
+-0 --> rewritten as a serial loop
+-1 --> rewritten as a serial loop

----- Before Optimisation -----

Parallel region 0:
+-2 (parallel)
+-0 (parallel)
+-1 (parallel)

----- After Optimisation -----

Parallel region 0:
+-2 (parallel)
+-0 (serial)
+-1 (serial)

Parallel region 0 (loop #2) had 0 loop(s) fused and 2 loop(s) serialized as part of the larger parallel loop (#2).

----- Loop invariant code motion -----

Allocation hoisting:
The memory allocation derived from the instruction at
/Users/qizhixuan/Library/CloudStorage/OneDrive-Personal/2024
Fall/MLE/workspace/mod3-Navxihziq/minitorch/fast_ops.py (182) is hoisted out of
the parallel loop labelled #2 (it will be performed before the loop is executed
and reused inside the loop):
Allocation:: out_index = np.zeros(len(out_shape), dtype=np.int32) # buffer
- numpy.empty() is used for the allocation.

The memory allocation derived from the instruction at
/Users/qizhixuan/Library/CloudStorage/OneDrive-Personal/2024
Fall/MLE/workspace/mod3-Navxihziq/minitorch/fast_ops.py (183) is hoisted out of
the parallel loop labelled #2 (it will be performed before the loop is executed
and reused inside the loop):
Allocation:: in_index = np.zeros(len(in_shape), dtype=np.int32) # buffer
- numpy.empty() is used for the allocation.

None
ZIP

```
=====
Parallel Accelerator Optimizing: Function tensor_zip.<locals>._zip,
/Users/qizhixuan/Library/CloudStorage/OneDrive-Personal/2024
Fall/MLE/workspace/mod3-Navxihziq/minitorch/fast_ops.py (214)
=====
```

```
Parallel loop listing for Function tensor_zip.<locals>._zip, /Users/qizhixuan/Library/CloudStorage/OneDrive-Personal/2024 Fall/MLE/workspace/mod3-Navxihziq/minitorch/fast_ops.py (214)
```

```
----- loop #10 -----
def _zip(
    out: Storage,
    out_shape: Shape,
    out_strides: Strides,
    a_storage: Storage,
    a_shape: Shape,
    a_strides: Strides,
    b_storage: Storage,
    b_shape: Shape,
    b_strides: Strides,
) -> None:
    # TODO: Implement for Task 3.1.
    # coerce the shape to int32
    out_shape = out_shape.astype(np.int32)
```

```

a_shape = a_shape.astype(np.int32)
b_shape = b_shape.astype(np.int32)
# TODO: check if out, a, b are stride-aligned
# if (
#     len(out_shape) == len(a_shape) == len(b_shape)
#     and np.array_equal(out_shape, a_shape)
#     and np.array_equal(out_shape, b_shape)
#     and np.array_equal(out_strides, a_strides)
#     and np.array_equal(out_strides, b_strides)
# ):
#     for i in prange(len(out)):
#         out[i] = fn(a_storage[i], b_storage[i])
# else:
for i in prange(len(out)):-----#6
    out_index = np.zeros(len(out_shape), dtype=np.int32) # buffer-----#3
    a_index = np.zeros(len(a_shape), dtype=np.int32) # buffer-----#4
    b_index = np.zeros(len(b_shape), dtype=np.int32) # buffer-----#5
    to_index(i, out_shape, out_index)
    broadcast_index(out_index, out_shape, a_shape, a_index)
    broadcast_index(out_index, out_shape, b_shape, b_index)

    out[i] = fn(
        a_storage[index_to_position(a_index, a_strides)],
        b_storage[index_to_position(b_index, b_strides)],
    )

```

----- Fusing loops -----
Attempting fusion of parallel loops (combines loops with similar properties)...
Following the attempted fusion of parallel for-loops there are 4 parallel for-loops (originating from loops labelled: #6, #3, #4, #5).

----- Optimising loop nests -----
Attempting loop nest rewrites (optimising for the largest parallel loops)...

←6 is a parallel loop
→3 → rewritten as a serial loop
→4 → rewritten as a serial loop
→5 → rewritten as a serial loop

----- Before Optimisation -----
Parallel region 0:
←6 (parallel)
→3 (parallel)
→4 (parallel)
→5 (parallel)

----- After Optimisation -----
Parallel region 0:
←6 (parallel)
→3 (serial)
→4 (serial)
→5 (serial)

Parallel region 0 (loop #6) had 0 loop(s) fused and 3 loop(s) serialized as part of the larger parallel loop (#6).

----- Loop invariant code motion -----
Allocation hoisting:
The memory allocation derived from the instruction at
/Users/qizhixuan/Library/CloudStorage/OneDrive-Personal/2024
Fall/MLE/workspace/mod3-Navxihziq/minitorch/fast_ops.py (242) is hoisted out of
the parallel loop labelled #6 (it will be performed before the loop is executed
and reused inside the loop):
Allocation:: out_index = np.zeros(len(out_shape), dtype=np.int32) # buffer
- numpy.empty() is used for the allocation.
The memory allocation derived from the instruction at
/Users/qizhixuan/Library/CloudStorage/OneDrive-Personal/2024
Fall/MLE/workspace/mod3-Navxihziq/minitorch/fast_ops.py (243) is hoisted out of
the parallel loop labelled #6 (it will be performed before the loop is executed
and reused inside the loop):
Allocation:: a_index = np.zeros(len(a_shape), dtype=np.int32) # buffer
- numpy.empty() is used for the allocation.
The memory allocation derived from the instruction at
/Users/qizhixuan/Library/CloudStorage/OneDrive-Personal/2024
Fall/MLE/workspace/mod3-Navxihziq/minitorch/fast_ops.py (244) is hoisted out of
the parallel loop labelled #6 (it will be performed before the loop is executed
and reused inside the loop):
Allocation:: b_index = np.zeros(len(b_shape), dtype=np.int32) # buffer
- numpy.empty() is used for the allocation.

None
REDUCE

===== Parallel Accelerator Optimizing: Function tensor_reduce.<locals>._reduce, /Users/qizhixuan/Library/CloudStorage/OneDrive-Personal/2024 Fall/MLE/workspace/mod3-Navxihziq/minitorch/fast_ops.py (279) =====

Parallel loop listing for Function tensor_reduce.<locals>._reduce, /Users/qizhixuan/Library/CloudStorage/OneDrive-Personal/2024 Fall/MLE/workspace/mod3-Navxihziq/minitorch/fast_ops.py (279)

```

def _reduce(
    out: Storage,
    out_shape: Shape,
    out_strides: Strides,
    a_storage: Storage,
    a_shape: Shape,
    a_strides: Strides,
    reduce_dim: int,
) -> None:
    # TODO: Implement for Task 3.1.
    # coerce the shape to int32
    out_shape = out_shape.astype(np.int32)
    a_shape = a_shape.astype(np.int32)
    for i in prange(len(out)):-----#9
        out_index = np.zeros(len(out_shape), dtype=np.int32) # buffer-----#7
        a_index = np.zeros(len(a_shape), dtype=np.int32) # buffer-----#8
        to_index(i, out_shape, out_index)
        # copy the out_index to the a_index (except for the reduce dim)
        for j in range(len(a_shape) - 1):
            j = j if j < reduce_dim else j + 1

```

```

        a_index[j] = out_index[j]
        a_index[reduce_dim] = 0
        a_pos = index_to_position(a_index, a_strides)
        temp: float = a_storage[a_pos] # avoid inner access to global variable
        for j in range(1, a_shape[reduce_dim]):
            temp = fn(temp, float(a_storage[a_pos + j * a_strides[reduce_dim]]))
        out[i] = temp

```

----- Fusing loops -----
Attempting fusion of parallel loops (combines loops with similar properties)...
Following the attempted fusion of parallel for-loops there are 3 parallel for-loops (originating from loops labelled: #9, #7, #8).

----- Optimising loop nests -----
Attempting loop nest rewrites (optimising for the largest parallel loops)...

```

+--9 is a parallel loop
+--8 --> rewritten as a serial loop
+--7 --> rewritten as a serial loop

```

----- Before Optimisation -----
Parallel region 0:
+--9 (parallel)
+--8 (parallel)
+--7 (parallel)

----- After Optimisation -----
Parallel region 0:
+--9 (parallel)
+--8 (serial)
+--7 (serial)

Parallel region 0 (loop #9) had 0 loop(s) fused and 2 loop(s) serialized as part of the larger parallel loop (#9).

----- Loop invariant code motion -----
Allocation hoisting:
The memory allocation derived from the instruction at
/Users/qizhixuan/Library/CloudStorage/OneDrive-Personal/2024
Fall/MLE/workspace/mod3-Navxihziq/minitorch/fast_ops.py (293) is hoisted out of
the parallel loop labelled #9 (it will be performed before the loop is executed
and reused inside the loop):
Allocation:: out_index = np.zeros(len(out_shape), dtype=np.int32) # buffer
- numpy.empty() is used for the allocation.
The memory allocation derived from the instruction at
/Users/qizhixuan/Library/CloudStorage/OneDrive-Personal/2024
Fall/MLE/workspace/mod3-Navxihziq/minitorch/fast_ops.py (294) is hoisted out of
the parallel loop labelled #9 (it will be performed before the loop is executed
and reused inside the loop):
Allocation:: a_index = np.zeros(len(a_shape), dtype=np.int32) # buffer
- numpy.empty() is used for the allocation.

None
MATRIX MULTIPLY

=====

Parallel Accelerator Optimizing: Function _tensor_matrix_multiply,
/Users/qizhixuan/Library/CloudStorage/OneDrive-Personal/2024
Fall/MLE/workspace/mod3-Navxihziq/minitorch/fast_ops.py (311)

=====

Parallel loop listing for Function _tensor_matrix_multiply, /Users/qizhixuan/Library/CloudStorage/OneDrive-Personal/2024 Fall/MLE/workspace/mod3-Navxihziq/minitorch/fast_ops.py (311)

	loop #ID
def _tensor_matrix_multiply(out: Storage, out_shape: Shape, out_strides: Strides, a_storage: Storage, a_shape: Shape, a_strides: Strides, b_storage: Storage, b_shape: Shape, b_strides: Strides,) -> None: """NUMBA tensor matrix multiply function. Should work for any tensor shapes that broadcast as long as ... assert a_shape[-1] == b_shape[-2] ... Optimizations: * Outer loop in parallel * No index buffers or function calls * Inner loop should have no global writes, 1 multiply. Args: ----- out (Storage): storage for 'out' tensor out_shape (Shape): shape for 'out' tensor out_strides (Strides): strides for 'out' tensor a_storage (Storage): storage for 'a' tensor a_shape (Shape): shape for 'a' tensor a_strides (Strides): strides for 'a' tensor b_storage (Storage): storage for 'b' tensor b_shape (Shape): shape for 'b' tensor b_strides (Strides): strides for 'b' tensor Returns: ----- None : Fills in 'out' """ a_batch_stride = a_strides[0] if a_shape[0] > 1 else 0 b_batch_stride = b_strides[0] if b_shape[0] > 1 else 0 # TODO: Implement for Task 3.2.	

```

for i in prange(len(out)):------| #10
    # disassemble the index
    out_batch = i // (out_shape[-2] * out_shape[-1])
    out_j = (i % out_strides[0]) % out_shape[-1]
    out_i = (i % out_strides[0]) // out_shape[-1]

    a_pos = out_batch * a_batch_stride + out_i * a_strides[-2] + 0
    b_pos = out_batch * b_batch_stride + 0 + out_j * b_strides[-1]

    acc = 0.0
    for j in range(a_shape[-1]): # iterate along the shared dim
        a_location = a_pos + j * a_strides[-1]
        b_location = b_pos + j * b_strides[-2]
        acc += a_storage[a_location] * b_storage[b_location]
    out[i] = acc
----- Fusing loops -----
Attempting fusion of parallel loops (combines loops with similar properties)...
Following the attempted fusion of parallel for-loops there are 1 parallel for-
loop(s) (originating from loops labelled: #10).
-----
----- Before Optimisation -----
----- After Optimisation -----
Parallel structure is already optimal.
-----
----- Loop invariant code motion -----
Allocation hoisting:
No allocation hoisting found
None
~\OneDrive\2024 Fall\MLE\workspace\mod3-Navxihziq master +1

```