

```

Last login: Mon Nov 18 19:35:48 on ttys009
> cd OneDrive/2024/ Fall/MLE/workspace/mod3-Navxihziq
> python project/parallel_check.py
Traceback (most recent call last):
  File "/Users/qizhixuan/Library/CloudStorage/OneDrive-Personal/2024 Fall/MLE/workspace/mod3-Navxihziq/project/parallel_check.py", line 1, in <module>
    from numba import njit
ModuleNotFoundError: No module named 'numba'
> source .venv/bin/activate
> python project/parallel_check.py
MAP
OMP: Info #276: omp_set_nested routine deprecated, please use omp_set_max_active_levels instead.

```

```

=====
Parallel Accelerator Optimizing: Function tensor_map.<locals>._map,
/Users/qizhixuan/Library/CloudStorage/OneDrive-Personal/2024
Fall/MLE/workspace/mod3-Navxihziq/minitorch/fast_ops.py (163)
=====

```

Parallel loop listing for Function tensor\_map.<locals>.\_map, /Users/qizhixuan/Library/CloudStorage/OneDrive-Personal/2024 Fall/MLE/workspace/mod3-Navxihziq/minitorch/fast\_ops.py (163)

```

def _map(                                     | loop #ID
    out: Storage,                             |
    out_shape: Shape,                         |
    out_strides: Strides,                     |
    in_storage: Storage,                      |
    in_shape: Shape,                          |
    in_strides: Strides,                      |
) -> None:                                    |
    # TODO: Implement for Task 3.1.            |
    # check if out, in are stride-aligned     |
    # if out_strides == in_strides:           |
    #     for i in prange(len(out)):           |
    #         out[i] = fn(in_storage[i])       |
    # else:                                    |
    #     # TODO: check if out, in are stride-aligned |
    for i in prange(len(out)):                 | #2
        out_index = np.zeros(len(out_shape)) # buffer----- | #0
        in_index = np.zeros(len(in_shape)) # buffer----- | #1
        to_index(i, out_shape, out_index)
        broadcast_index(out_index, out_shape, in_shape, in_index)
        out[i] = fn(in_storage[index_to_position(in_index, in_strides)])
    ----- Fusing loops

```

Attempting fusion of parallel loops (combines loops with similar properties)...  
Following the attempted fusion of parallel for-loops there are 3 parallel for-loop(s) (originating from loops labelled: #2, #0, #1).

----- Optimising loop nests -----  
Attempting loop nest rewrites (optimising for the largest parallel loops)...

```

+--2 is a parallel loop
+--0 --> rewritten as a serial loop
+--1 --> rewritten as a serial loop

```

----- Before Optimisation -----

```

Parallel region 0:
+--2 (parallel)
+--0 (parallel)
+--1 (parallel)

```

----- After Optimisation -----

```

Parallel region 0:
+--2 (parallel)
+--0 (serial)
+--1 (serial)

```

Parallel region 0 (loop #2) had 0 loop(s) fused and 2 loop(s) serialized as part of the larger parallel loop (#2).

----- Loop invariant code motion -----

Allocation hoisting:  
The memory allocation derived from the instruction at  
/Users/qizhixuan/Library/CloudStorage/OneDrive-Personal/2024  
Fall/MLE/workspace/mod3-Navxihziq/minitorch/fast\_ops.py (179) is hoisted out of  
the parallel loop labelled #2 (it will be performed before the loop is executed  
and reused inside the loop):

```

Allocation:: out_index = np.zeros(len(out_shape)) # buffer
- numpy.empty() is used for the allocation.

```

The memory allocation derived from the instruction at  
/Users/qizhixuan/Library/CloudStorage/OneDrive-Personal/2024  
Fall/MLE/workspace/mod3-Navxihziq/minitorch/fast\_ops.py (180) is hoisted out of  
the parallel loop labelled #2 (it will be performed before the loop is executed  
and reused inside the loop):

```

Allocation:: in_index = np.zeros(len(in_shape)) # buffer
- numpy.empty() is used for the allocation.

```

None  
ZIP

```

=====
Parallel Accelerator Optimizing: Function tensor_zip.<locals>._zip,
/Users/qizhixuan/Library/CloudStorage/OneDrive-Personal/2024
Fall/MLE/workspace/mod3-Navxihziq/minitorch/fast_ops.py (210)
=====

```

Parallel loop listing for Function tensor\_zip.<locals>.\_zip, /Users/qizhixuan/Library/CloudStorage/OneDrive-Personal/2024 Fall/MLE/workspace/mod3-Navxihziq/minitorch/fast\_ops.py (210)

```

def _zip(                                     | loop #ID
    out: Storage,                             |
    out_shape: Shape,                         |
    out_strides: Strides,                     |
    a_storage: Storage,                       |
    a_shape: Shape,                           |
    a_strides: Strides,                       |
    b_storage: Storage,                       |
    b_shape: Shape,                           |
    b_strides: Strides,                       |
) -> None:                                    |
    # TODO: Implement for Task 3.1.            |
    # TODO: check if out, a, b are stride-aligned |

```

```

# if (
#     len(out_shape) == len(a_shape) == len(b_shape)
#     and np.array_equal(out_shape, a_shape)
#     and np.array_equal(out_shape, b_shape)
#     and np.array_equal(out_strides, a_strides)
#     and np.array_equal(out_strides, b_strides)
# ):
#     for i in prange(len(out)):
#         out[i] = fn(a_storage[i], b_storage[i])
# else:
#     for i in prange(len(out)):-----#6
#         out_index = np.zeros(len(out_shape)) # buffer-----#3
#         a_index = np.zeros(len(a_shape)) # buffer-----#4
#         b_index = np.zeros(len(b_shape)) # buffer-----#5
#         to_index(i, out_shape, out_index)
#         broadcast_index(out_index, out_shape, a_shape, a_index)
#         broadcast_index(out_index, out_shape, b_shape, b_index)
#
#         out[i] = fn(a_storage[index_to_position(a_index, a_strides)],
#                   b_storage[index_to_position(b_index, b_strides)])
----- Fusing loops
Attempting fusion of parallel loops (combines loops with similar properties)...
Following the attempted fusion of parallel for-loops there are 4 parallel for-
loop(s) (originating from loops labelled: #6, #3, #4, #5).

----- Optimising loop nests -----
Attempting loop nest rewrites (optimising for the largest parallel loops)...

+--6 is a parallel loop
+--3 --> rewritten as a serial loop
+--4 --> rewritten as a serial loop
+--5 --> rewritten as a serial loop
-----
----- Before Optimisation -----
Parallel region 0:
+--6 (parallel)
+--3 (parallel)
+--4 (parallel)
+--5 (parallel)

----- After Optimisation -----
Parallel region 0:
+--6 (parallel)
+--3 (serial)
+--4 (serial)
+--5 (serial)

Parallel region 0 (loop #6) had 0 loop(s) fused and 3 loop(s) serialized as part
of the larger parallel loop (#6).

----- Loop invariant code motion -----
Allocation hoisting:
The memory allocation derived from the instruction at
/Users/qizhixuan/Library/CloudStorage/OneDrive-Personal/2024
Fall/MLE/workspace/mod3-Navxihziq/minitorch/fast_ops.py (234) is hoisted out of
the parallel loop labelled #6 (it will be performed before the loop is executed
and reused inside the loop):
Allocation:: out_index = np.zeros(len(out_shape)) # buffer
- numpy.empty() is used for the allocation.
The memory allocation derived from the instruction at
/Users/qizhixuan/Library/CloudStorage/OneDrive-Personal/2024
Fall/MLE/workspace/mod3-Navxihziq/minitorch/fast_ops.py (235) is hoisted out of
the parallel loop labelled #6 (it will be performed before the loop is executed
and reused inside the loop):
Allocation:: a_index = np.zeros(len(a_shape)) # buffer
- numpy.empty() is used for the allocation.
The memory allocation derived from the instruction at
/Users/qizhixuan/Library/CloudStorage/OneDrive-Personal/2024
Fall/MLE/workspace/mod3-Navxihziq/minitorch/fast_ops.py (236) is hoisted out of
the parallel loop labelled #6 (it will be performed before the loop is executed
and reused inside the loop):
Allocation:: b_index = np.zeros(len(b_shape)) # buffer
- numpy.empty() is used for the allocation.
None
REDUCE

=====
Parallel Accelerator Optimizing: Function tensor_reduce.<locals>._reduce,
/Users/qizhixuan/Library/CloudStorage/OneDrive-Personal/2024
Fall/MLE/workspace/mod3-Navxihziq/minitorch/fast_ops.py (269)
=====

Parallel loop listing for Function tensor_reduce.<locals>._reduce, /Users/qizhixuan/Library/CloudStorage/OneDrive-Personal/2024 Fall/MLE/workspace/mod3-Navxihziq/minitorch/fast_ops.py (269)
-----|loop #ID
def _reduce(
    out: Storage,
    out_shape: Shape,
    out_strides: Strides,
    a_storage: Storage,
    a_shape: Shape,
    a_strides: Strides,
    reduce_dim: int,
) -> None:
    # TODO: Implement for Task 3.1.
    for i in prange(len(out)):-----#9
        out_index = np.zeros(len(out_shape)) # buffer-----#7
        a_index = np.zeros(len(a_shape)) # buffer-----#8
        to_index(i, out_shape, out_index)
        # copy the out_index to the a_index (except for the reduce dim)
        for j in range(len(a_shape)-1):
            j = j if j < reduce_dim else j + 1
            a_index[j] = out_index[j]

        a_index[reduce_dim] = 0
        a_pos = index_to_position(a_index, a_strides)
        temp = a_storage[a_pos] # avoid inner access to global variable
        for j in range(1, a_shape[reduce_dim]):
            temp = fn(temp, a_storage[a_pos + j * a_strides[reduce_dim]])
        out[i] = temp

```

```

----- Fusing loops -----
Attempting fusion of parallel loops (combines loops with similar properties)...
Following the attempted fusion of parallel for-loops there are 3 parallel for-
loop(s) (originating from loops labelled: #9, #7, #8).

----- Optimising loop nests -----
Attempting loop nest rewrites (optimising for the largest parallel loops)...

+-9 is a parallel loop
+-8 --> rewritten as a serial loop
+-7 --> rewritten as a serial loop

----- Before Optimisation -----
Parallel region 0:
+-9 (parallel)
+-8 (parallel)
+-7 (parallel)

----- After Optimisation -----
Parallel region 0:
+-9 (parallel)
+-8 (serial)
+-7 (serial)

Parallel region 0 (loop #9) had 0 loop(s) fused and 2 loop(s) serialized as part
of the larger parallel loop (#9).

----- Loop invariant code motion -----
Allocation hoisting:
The memory allocation derived from the instruction at
/Users/qizhixuan/Library/CloudStorage/OneDrive-Personal/2024
Fall/MLE/workspace/mod3-Navxihziq/minitorch/fast_ops.py (280) is hoisted out of
the parallel loop labelled #9 (it will be performed before the loop is executed
and reused inside the loop):
Allocation:: out_index = np.zeros(len(out_shape)) # buffer
- numpy.empty() is used for the allocation.
The memory allocation derived from the instruction at
/Users/qizhixuan/Library/CloudStorage/OneDrive-Personal/2024
Fall/MLE/workspace/mod3-Navxihziq/minitorch/fast_ops.py (281) is hoisted out of
the parallel loop labelled #9 (it will be performed before the loop is executed
and reused inside the loop):
Allocation:: a_index = np.zeros(len(a_shape)) # buffer
- numpy.empty() is used for the allocation.
None
MATRIX MULTIPLY

=====
Parallel Accelerator Optimizing: Function _tensor_matrix_multiply,
/Users/qizhixuan/Library/CloudStorage/OneDrive-Personal/2024
Fall/MLE/workspace/mod3-Navxihziq/minitorch/fast_ops.py (297)
=====

```

Parallel loop listing for Function \_tensor\_matrix\_multiply, /Users/qizhixuan/Library/CloudStorage/OneDrive-Personal/2024 Fall/MLE/workspace/mod3-Navxihziq/minitorch/fast\_ops.py (297)

```

def _tensor_matrix_multiply(
    out: Storage,
    out_shape: Shape,
    out_strides: Strides,
    a_storage: Storage,
    a_shape: Shape,
    a_strides: Strides,
    b_storage: Storage,
    b_shape: Shape,
    b_strides: Strides,
) -> None:
    """NUMBA tensor matrix multiply function.

    Should work for any tensor shapes that broadcast as long as
    ...
    assert a_shape[-1] == b_shape[-2]
    ...

    Optimizations:

    * Outer loop in parallel
    * No index buffers or function calls
    * Inner loop should have no global writes, 1 multiply.

    Args:
    -----
        out (Storage): storage for 'out' tensor
        out_shape (Shape): shape for 'out' tensor
        out_strides (Strides): strides for 'out' tensor
        a_storage (Storage): storage for 'a' tensor
        a_shape (Shape): shape for 'a' tensor
        a_strides (Strides): strides for 'a' tensor
        b_storage (Storage): storage for 'b' tensor
        b_shape (Shape): shape for 'b' tensor
        b_strides (Strides): strides for 'b' tensor

    Returns:
    -----
        None : Fills in 'out'

    """
    a_batch_stride = a_strides[0] if a_shape[0] > 1 else 0
    b_batch_stride = b_strides[0] if b_shape[0] > 1 else 0

    # TODO: Implement for Task 3.2.
    for i in prange(len(out)):
        # disassemble the index
        out_batch = i // (out_shape[-2] * out_shape[-1])
        out_i = (i // out_shape[-1]) % out_shape[-2]
        out_j = i % out_shape[-1]

        a_pos = out_batch * a_batch_stride + out_i * a_strides[-2] + 0
        b_pos = out_batch * b_batch_stride + 0 + out_j * b_strides[-1]

```

```
acc = 0.0
for j in range(a_shape[-1]): # iterate along the shared dim
    acc += (a_storage[a_pos + j * a_strides[-1]] *
            b_storage[b_pos + j * b_strides[-2]])
out[i] = acc
----- Fusing loops -----
Attempting fusion of parallel loops (combines loops with similar properties)...
Following the attempted fusion of parallel for-loops there are 1 parallel for-
loop(s) (originating from loops labelled: #10).
----- Before Optimisation -----
----- After Optimisation -----
Parallel structure is already optimal.
----- Loop invariant code motion -----
Allocation hoisting:
No allocation hoisting found
None
🍏 ~/OneDrive/2024 Fall/MLE/workspace/mod3-Navxihziq 🍏 master !3
```