

```
Last login: Thu Nov 21 22:55:41 on ttys010
> ssh m1e
Welcome to Ubuntu 22.04.5 LTS (GNU/Linux 6.8.0-48-generic x86_64)
```



```
* Documentation:  https://help.ubuntu.com
* Management:    https://landscape.canonical.com
* Support:        https://ubuntu.com/pro
```

```
Last login: Fri Nov 22 03:55:49 2024 from 52.124.35.169
```

```
ubuntu@146-235-212-176:~$ cd proj/mod3-Navxihziq/
(.venv) ubuntu@146-235-212-176:~/proj/mod3-Navxihziq$ source .venv/bin/activate
(.venv) ubuntu@146-235-212-176:~/proj/mod3-Navxihziq$ python project/run_fast_tensor.py --BACKEND gpu --HIDDEN 100 --DATASET xor --RATE 0.05
/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/dispatcher.py:536: NumbaPerformanceWarning: Grid size 1 will likely result in GPU under-utilization due to low occupancy.
  warn(NumbaPerformanceWarning(msg))
/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/cudadrv/devicearray.py:888: NumbaPerformanceWarning: Host array used in CUDA kernel will incur copy overhead to/from device.
  warn(NumbaPerformanceWarning(msg))
/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/dispatcher.py:536: NumbaPerformanceWarning: Grid size 7 will likely result in GPU under-utilization due to low occupancy.
  warn(NumbaPerformanceWarning(msg))
/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/cudadrv/devicearray.py:888: NumbaPerformanceWarning: Host array used in CUDA kernel will incur copy overhead to/from device.
  warn(NumbaPerformanceWarning(msg))
/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/dispatcher.py:536: NumbaPerformanceWarning: Grid size 4 will likely result in GPU under-utilization due to low occupancy.
  warn(NumbaPerformanceWarning(msg))
/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/dispatcher.py:536: NumbaPerformanceWarning: Grid size 1 will likely result in GPU under-utilization due to low occupancy.
  warn(NumbaPerformanceWarning(msg))
/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/cudadrv/devicearray.py:888: NumbaPerformanceWarning: Host array used in CUDA kernel will incur copy overhead to/from device.
  warn(NumbaPerformanceWarning(msg))
/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/dispatcher.py:536: NumbaPerformanceWarning: Grid size 7 will likely result in GPU under-utilization due to low occupancy.
  warn(NumbaPerformanceWarning(msg))
/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/dispatcher.py:536: NumbaPerformanceWarning: Grid size 4 will likely result in GPU under-utilization due to low occupancy.
  warn(NumbaPerformanceWarning(msg))
/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/cudadrv/devicearray.py:888: NumbaPerformanceWarning: Host array used in CUDA kernel will incur copy overhead to/from device.
  warn(NumbaPerformanceWarning(msg))
/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/dispatcher.py:536: NumbaPerformanceWarning: Grid size 32 will likely result in GPU under-utilization due to low occupancy.
  warn(NumbaPerformanceWarning(msg))
/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/cudadrv/devicearray.py:888: NumbaPerformanceWarning: Host array used in CUDA kernel will incur copy overhead to/from device.
  warn(NumbaPerformanceWarning(msg))
/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/dispatcher.py:536: NumbaPerformanceWarning: Grid size 32 will likely result in GPU under-utilization due to low occupancy.
  warn(NumbaPerformanceWarning(msg))
/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/dispatcher.py:536: NumbaPerformanceWarning: Grid size 4 will likely result in GPU under-utilization due to low occupancy.
  warn(NumbaPerformanceWarning(msg))
/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/dispatcher.py:536: NumbaPerformanceWarning: Grid size 1 will likely result in GPU under-utilization due to low occupancy.
  warn(NumbaPerformanceWarning(msg))
/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/cudadrv/devicearray.py:888: NumbaPerformanceWarning: Host array used in CUDA kernel will incur copy overhead to/from device.
  warn(NumbaPerformanceWarning(msg))
/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/dispatcher.py:536: NumbaPerformanceWarning: Grid size 1 will likely result in GPU under-utilization due to low occupancy.
  warn(NumbaPerformanceWarning(msg))
/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/cudadrv/devicearray.py:888: NumbaPerformanceWarning: Host array used in CUDA kernel will incur copy overhead to/from device.
  warn(NumbaPerformanceWarning(msg))
/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/dispatcher.py:536: NumbaPerformanceWarning: Grid size 1 will likely result in GPU under-utilization due to low occupancy.
  warn(NumbaPerformanceWarning(msg))
/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/cudadrv/devicearray.py:888: NumbaPerformanceWarning: Host array used in CUDA kernel will incur copy overhead to/from device.
  warn(NumbaPerformanceWarning(msg))
/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/dispatcher.py:536: NumbaPerformanceWarning: Grid size 1 will likely result in GPU under-utilization due to low occupancy.
  warn(NumbaPerformanceWarning(msg))
/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/cudadrv/devicearray.py:888: NumbaPerformanceWarning: Host array used in CUDA kernel will incur copy overhead to/from device.
  warn(NumbaPerformanceWarning(msg))
/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/dispatcher.py:536: NumbaPerformanceWarning: Grid size 4 will likely result in GPU under-utilization due to low occupancy.
  warn(NumbaPerformanceWarning(msg))
/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/dispatcher.py:536: NumbaPerformanceWarning: Grid size 32 will likely result in GPU under-utilization due to low occupancy.
  warn(NumbaPerformanceWarning(msg))
/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/cudadrv/devicearray.py:888: NumbaPerformanceWarning: Host array used in CUDA kernel will incur copy overhead to/from device.
  warn(NumbaPerformanceWarning(msg))
/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/dispatcher.py:536: NumbaPerformanceWarning: Grid size 1 will likely result in GPU under-utilization due to low occupancy.
  warn(NumbaPerformanceWarning(msg))
/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/cudadrv/devicearray.py:888: NumbaPerformanceWarning: Host array used in CUDA kernel will incur copy overhead to/from device.
  warn(NumbaPerformanceWarning(msg))
/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/dispatcher.py:536: NumbaPerformanceWarning: Grid size 100 will likely result in GPU under-utilization due to low occupancy.
  warn(NumbaPerformanceWarning(msg))
/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/dispatcher.py:536: NumbaPerformanceWarning: Grid size 16 will likely result in GPU under-utilization due to low occupancy.
  warn(NumbaPerformanceWarning(msg))
/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/dispatcher.py:536: NumbaPerformanceWarning: Grid size 7 will likely result in GPU under-utilization due to low occupancy.
  warn(NumbaPerformanceWarning(msg))
/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/dispatcher.py:536: NumbaPerformanceWarning: Grid size 4 will likely result in GPU under-utilization due to low occupancy.
  warn(NumbaPerformanceWarning(msg))
/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/dispatcher.py:536: NumbaPerformanceWarning: Grid size 8 will likely result in GPU under-utilization due to low occupancy.
  warn(NumbaPerformanceWarning(msg))
/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/dispatcher.py:536: NumbaPerformanceWarning: Grid size 2 will likely result in GPU under-utilization due to low o
```

```
ccupancy.
warn(NumbaPerformanceWarning(msg))
/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/cudadrv/devicearray.py:888: NumbaPerformanceWarning: Host array used in CUDA kernel will incur copy overhead to/from device.
warn(NumbaPerformanceWarning(msg))
/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/dispatcher.py:536: NumbaPerformanceWarning: Grid size 2 will likely result in GPU under-utilization due to low occupancy.
warn(NumbaPerformanceWarning(msg))
/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/cudadrv/devicearray.py:888: NumbaPerformanceWarning: Host array used in CUDA kernel will incur copy overhead to/from device.
warn(NumbaPerformanceWarning(msg))
/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/dispatcher.py:536: NumbaPerformanceWarning: Grid size 2 will likely result in GPU under-utilization due to low occupancy.
warn(NumbaPerformanceWarning(msg))
Epoch 0 loss 7.929135135737758 correct 29 time 3.79s time per epoch 3.79s
Epoch 10 loss 4.552853840425255 correct 39 time 18.11s time per epoch 1.65s
Epoch 20 loss 4.03626367462243 correct 46 time 32.44s time per epoch 1.54s
Epoch 30 loss 5.893117552297534 correct 42 time 46.35s time per epoch 1.50s
Epoch 40 loss 2.4576929605961153 correct 46 time 60.18s time per epoch 1.47s
Epoch 50 loss 1.535034610590493 correct 48 time 74.00s time per epoch 1.45s
Epoch 60 loss 3.2892836280156157 correct 48 time 80.12s time per epoch 1.44s
Epoch 70 loss 3.251929496623473 correct 46 time 101.94s time per epoch 1.44s
Epoch 80 loss 1.9497109010185258 correct 48 time 115.79s time per epoch 1.43s
Epoch 90 loss 1.5444750783124537 correct 48 time 129.63s time per epoch 1.42s
Epoch 100 loss 1.1125971672677224 correct 49 time 143.46s time per epoch 1.42s
Epoch 110 loss 3.682435553173611 correct 46 time 157.32s time per epoch 1.42s
Epoch 120 loss 1.7617037055094134 correct 49 time 171.43s time per epoch 1.42s
Epoch 130 loss 1.973978643164274 correct 48 time 186.53s time per epoch 1.42s
Epoch 140 loss 0.6202532269373142 correct 49 time 201.80s time per epoch 1.43s
Epoch 150 loss 1.1928340105537187 correct 49 time 217.35s time per epoch 1.44s
Epoch 160 loss 0.4595514254053154 correct 49 time 232.68s time per epoch 1.45s
Epoch 170 loss 0.6903858274016108 correct 49 time 247.76s time per epoch 1.45s
Epoch 180 loss 0.4490669854618037 correct 49 time 262.90s time per epoch 1.45s
Epoch 190 loss 1.2440319869789267 correct 50 time 278.50s time per epoch 1.46s
Epoch 200 loss 0.3069011765969647 correct 49 time 293.75s time per epoch 1.46s
Epoch 210 loss 0.8205602489490569 correct 50 time 309.02s time per epoch 1.46s
Epoch 220 loss 0.760351541943433 correct 50 time 324.18s time per epoch 1.47s
Epoch 230 loss 0.20266915345083689 correct 50 time 339.34s time per epoch 1.47s
Epoch 240 loss 0.8837380207213554 correct 50 time 354.51s time per epoch 1.47s
Epoch 250 loss 0.6624675861406382 correct 50 time 369.80s time per epoch 1.47s
Epoch 260 loss 0.27289019807549864 correct 50 time 385.29s time per epoch 1.48s
Epoch 270 loss 0.4461875158112574 correct 50 time 400.77s time per epoch 1.48s
Epoch 280 loss 0.094890440850543254 correct 50 time 416.33s time per epoch 1.48s
Epoch 290 loss 0.4086393045778214 correct 50 time 431.48s time per epoch 1.48s
Epoch 300 loss 0.18219400200469532 correct 50 time 446.75s time per epoch 1.48s
Epoch 310 loss 0.5106873530330133 correct 50 time 462.13s time per epoch 1.49s
Epoch 320 loss 0.45995888113596783 correct 50 time 477.67s time per epoch 1.49s
Epoch 330 loss 0.1200706412667914 correct 50 time 493.10s time per epoch 1.49s
Epoch 340 loss 0.23290781846236983 correct 50 time 508.35s time per epoch 1.49s
Epoch 350 loss 0.20696767158308954 correct 50 time 523.53s time per epoch 1.49s
Epoch 360 loss 0.24870764326832792 correct 50 time 538.91s time per epoch 1.49s
Epoch 370 loss 0.27658419721677224 correct 50 time 554.12s time per epoch 1.49s
Epoch 380 loss 0.3101808040909413 correct 50 time 568.82s time per epoch 1.49s
Epoch 390 loss 0.17399080305812901 correct 50 time 582.94s time per epoch 1.49s
Epoch 400 loss 0.25045310158620887 correct 50 time 596.92s time per epoch 1.49s
^CTraceback (most recent call last):
  File "/home/ubuntu/proj/mod3-Navxihziq/project/run_fast_tensor.py", line 145, in <module>
    ).train(data, RATE)
  File "/home/ubuntu/proj/mod3-Navxihziq/project/run_fast_tensor.py", line 100, in train
    (loss / y.shape[0]).sum().view(1).backward()
  File "/home/ubuntu/proj/mod3-Navxihziq/minitorch/tensor.py", line 278, in backward
    backpropagate(self, grad_output)
  File "/home/ubuntu/proj/mod3-Navxihziq/minitorch/autodiff.py", line 168, in backpropagate
    deriv_map[parent.unique_id] += deriv
  File "/home/ubuntu/proj/mod3-Navxihziq/minitorch/tensor.py", line 328, in __radd__
    return Add.apply(self._ensure_tensor(b), self)
  File "/home/ubuntu/proj/mod3-Navxihziq/minitorch/tensor_functions.py", line 54, in apply
    c = cls._forward(ctx, *raw_vals)
  File "/home/ubuntu/proj/mod3-Navxihziq/minitorch/tensor_functions.py", line 38, in _forward
    return cls._forward(ctx, *args) # type: ignore
  File "/home/ubuntu/proj/mod3-Navxihziq/minitorch/tensor_functions.py", line 91, in _forward
    return t1.f.add_zip(t1, t2)
  File "/home/ubuntu/proj/mod3-Navxihziq/minitorch/cuda_ops.py", line 78, in ret
    out = a.zeros(c_shape)
  File "/home/ubuntu/proj/mod3-Navxihziq/minitorch/tensor.py", line 214, in zeros
    out._type_(self.backend)
  File "/home/ubuntu/proj/mod3-Navxihziq/minitorch/tensor.py", line 150, in _type_
    self._tensor.to_cuda()
  File "/home/ubuntu/proj/mod3-Navxihziq/minitorch/tensor_data.py", line 218, in to_cuda_
    self._storage = numba.cuda.to_device(self._storage)
  File "/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/cudadrv/devices.py", line 232, in _require_cuda_context
    return fn(*args, **kws)
  File "/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/api.py", line 128, in to_device
    to, new = devicearray.auto_device(obj, stream=stream, copy=copy,
  File "/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/cudadrv/devicearray.py", line 889, in auto_device
    devobj.copy_to_device(obj, stream=stream)
  File "/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/cudadrv/devices.py", line 232, in _require_cuda_context
    return fn(*args, **kws)
  File "/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/cudadrv/devicearray.py", line 223, in copy_to_device
    sentry_contiguous(self)
  File "/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/cudadrv/devicearray.py", line 849, in sentry_contiguous
    core = array_core(ary)
  File "/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/cudadrv/devicearray.py", line 823, in array_core
    return ary[tuple(core_index)]
  File "/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/cudadrv/devices.py", line 231, in _require_cuda_context
    with runtime.ensure_context():
  File "/usr/lib/python3.10/contextlib.py", line 135, in __enter__
    return next(self.gen)
  File "/home/ubuntu/proj/mod3-Navxihziq/.venv/lib/python3.10/site-packages/numba/cuda/cudadrv/devices.py", line 122, in ensure_context
    oldctx = self._get_attached_context()
KeyboardInterrupt

(.venv) ubuntu@146-235-212-176:~/proj/mod3-Navxihziq$
```